


Lecture 4:

Temporal Difference Learning


B. Ravindran

Dynamic Programming

$$\underline{v_\pi(s)} = \sum_a \pi(a | s) \sum_{s', r} p(s', r | s, a) [r + \gamma \underline{v_\pi(s')}]$$




Simplest “TD” Method




Simplest “TD” Method

$$V(s_t) \leftarrow V(s_t) + \alpha [r_{t+1} + \gamma V(s_{t+1}) - V(s_t)]$$




Simplest “TD” Method

$$V(s_t) \leftarrow V(s_t) + \alpha [r_{t+1} + \gamma V(s_{t+1}) - V(s_t)]$$




Temporal Difference

- Simple rule to explain complex behaviors
 - Intuition: Prediction of outcome at time $t+1$ is better than the prediction at time t . Hence use the later prediction to adjust the earlier prediction.
 - Has had profound impact in behavioral psychology and neuroscience!
- 

Bootstrapping and Sampling

- Bootstrapping: Update using an estimate

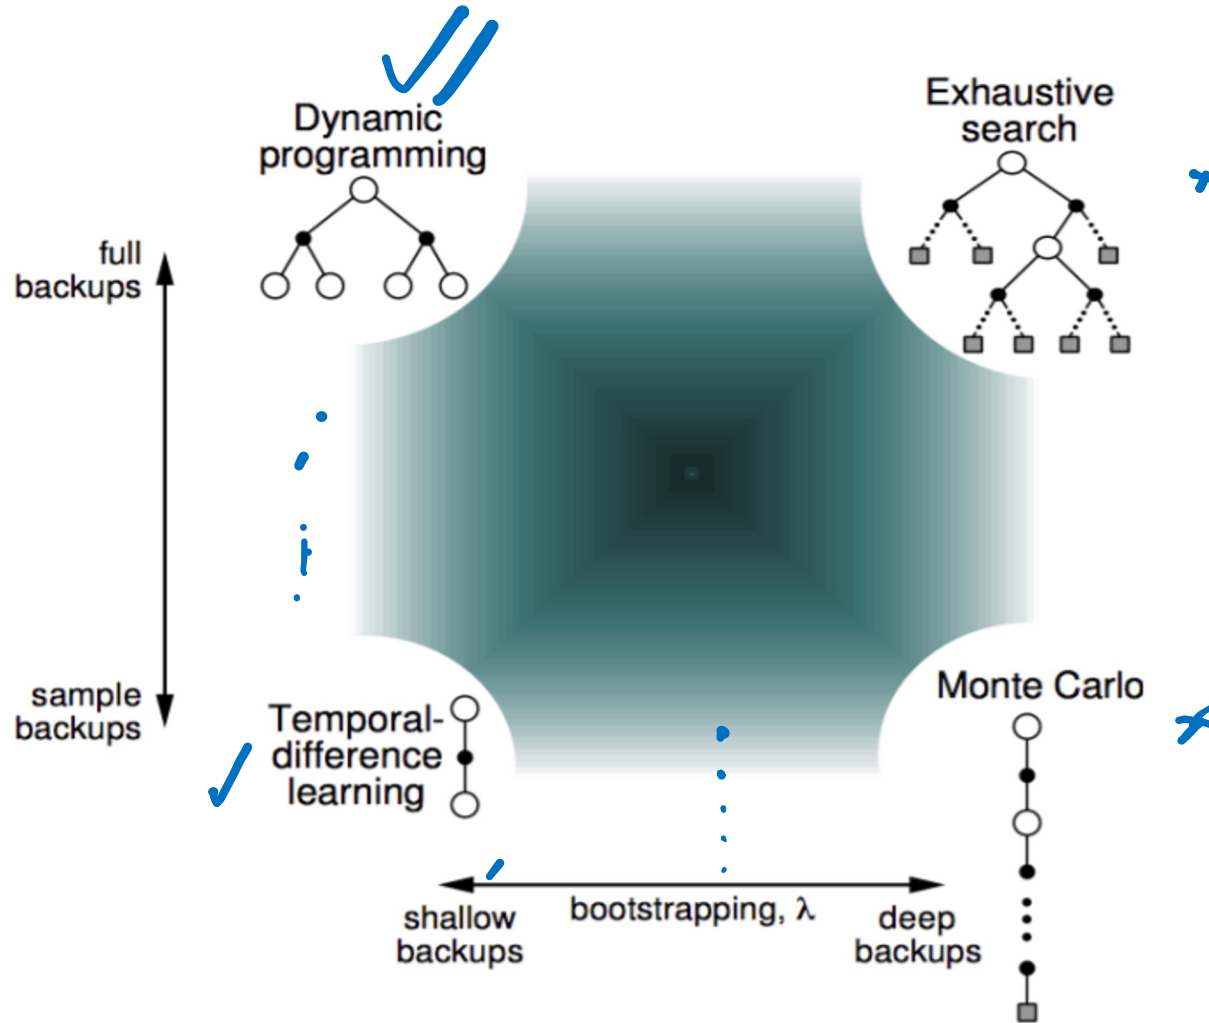
- DP and TD bootstrap
- Monte Carlo does not bootstrap.


$$V_{\pi}(s) = E_{\pi} \{ u_t | s_t = s \}$$

- Sampling: Update calculated using samples without model

- TD and Monte Carlo sample.
- DP (typically) does not sample

Bootstrapping and Sampling



Advantages of TD

S_t

A_t

$\sim S_{t+1} R_{t+1}$

- TD methods do not require a model of the environment, only experience (sampling)
- TD methods can be fully incremental (bootstrapping)
 - You can learn **before** knowing the final outcome
 - Less memory
 - Less peak computation.
 - You can learn **without** the final outcome
 - From incomplete sequences

Full Model
Distributional model.
 $p(s', r | s, a)$

Simulations
Sample Model.
 $s', r | s, a \sim p(s', r | s, a)$


TD Prediction

- Policy Evaluation (the prediction problem): for a given policy, compute the state-value function.
- No knowledge of p and r , but access to the real system, or a “sample” model assumed.
- Uses “bootstrapping” and sampling

π
 \mathcal{V}_π

The simplest TD method, TD(0):

$$V_{k+1}(s_t) \leftarrow V_k(s_t) + \alpha[r_{t+1} + \gamma V_k(s_{t+1}) - V_k(s_t)]$$



Stochastic Averaging Rule

$$E(x) \approx \frac{1}{n} \sum_{i=1}^n x_i$$

Let \bar{x}_n be the average of the first n samples

$$\begin{aligned}\bar{x}_{n+1} &= \frac{1}{n+1} (x_{n+1} + n\bar{x}_n) \\ &= \frac{1}{n+1} (x_{n+1} + n\bar{x}_n + \bar{x}_n - \bar{x}_n) \\ &= \frac{1}{n+1} ((n+1)\bar{x}_n + (x_{n+1} - \bar{x}_n)) \\ &= \bar{x}_n + \frac{1}{n+1} (x_{n+1} - \bar{x}_n) \\ &= \bar{x}_n + \alpha (x_{n+1} - \bar{x}_n)\end{aligned}$$

new estimate = old estimate + α (new sample - old estimate)

Stochastic Averaging Rule

$$E(x) \approx \frac{1}{n} \sum_{i=1}^n x_i$$


Let \bar{x}_n be the average of the first n samples

$$\begin{aligned}\bar{x}_{n+1} &= \frac{1}{n+1} (x_{n+1} + n\bar{x}_n) \\ &= \frac{1}{n+1} (x_{n+1} + n\bar{x}_n + \bar{x}_n - \bar{x}_n) \\ &= \frac{1}{n+1} ((n+1)\bar{x}_n + (x_{n+1} - \bar{x}_n)) \\ &= \bar{x}_n + \frac{1}{n+1} (x_{n+1} - \bar{x}_n) \\ &= \bar{x}_n + \alpha (x_{n+1} - \bar{x}_n)\end{aligned}$$

new estimate = old estimate + α (new sample - old estimate)

$$V_{k+1}(s_t) \leftarrow V_k(s_t) + \alpha [r_{t+1} + \gamma V_k(s_{t+1}) - V_k(s_t)]$$

TD Update Example



Assuming :

reward, r , $A \rightarrow B : 0$

$\alpha : 0.2$


$\gamma : 0.9$

$$V(A) = V(A) + \alpha[r + \gamma V(B) - V(A)]$$

$$V(A) = 0.5 + 0.2[0 + 0.9 * 0.8 - 0.5] = 0.544$$

TD Control

- The control problem: approximate optimal policies.
- Recall the idea of GPI:



- Policy evaluation: use TD(0) to evaluate value function.
- Policy improvement: make policy **greedy** wrt current value function.
- Note that we estimate **action values** rather than **state values** in the absence of a model.

ε -Greedy Policies

$$a^* \leftarrow \arg \max_a Q(s, a)$$

For all $a \in \mathcal{A}(s)$:

$$\pi(a|s) \leftarrow \begin{cases} 1 - \varepsilon + \varepsilon/|\mathcal{A}(s)| & \text{if } a = a^* \\ \varepsilon/|\mathcal{A}(s)| & \text{if } a \neq a^* \end{cases}$$

- Any ε -greedy policy with respect to Q following π is an improvement over any ε -soft policy π is assured by the policy improvement theorem

Sarsa: On-Policy TD Control


- In on-policy control, we try improving the policy used for making decisions.

SARSA

After every transition from a nonterminal state s_t , do this:


$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha [r_{t+1} + \gamma Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t)]$$

If s_{t+1} is terminal, then $Q(s_{t+1}, a_{t+1}) = 0$.



- Convergence is guaranteed as long as
 - all state-action pairs are visited an infinite number of times
 - the policy converges in the limit to the greedy policy

Sarsa: On-Policy TD Control



$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha [r_{t+1} + \gamma Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t)]$$

Sarsa Algorithm

Initialize $Q(s,a)$ arbitrarily

Repeat (for each episode)

Initialize s

Choose a from s using policy derived from Q (e.g., ϵ - greedy)

Repeat (for each step of episode):

Take action a , observe r, s'

Choose a' from s' using policy derived from Q (e.g., ϵ - greedy)

$$Q(s,a) \leftarrow Q(s,a) + \alpha[r + \gamma Q(s',a') - Q(s,a)]$$

$$s \leftarrow s'; a \leftarrow a';$$

until s is terminal

Sarsa Algorithm

Initialize $Q(s,a)$ arbitrarily

Repeat (for each episode)

Initialize s

Choose a from s using policy derived from Q (e.g., ϵ - greedy)

Repeat (for each step of episode):

Take action a , observe r, s'


IMPROVEMENT Choose a' from s' using policy derived from Q (e.g., ϵ - greedy)

EVALUATION $Q(s,a) \leftarrow Q(s,a) + \alpha[r + \gamma Q(s',a') - Q(s,a)]$

$s \leftarrow s'; a \leftarrow a';$

until s is terminal

Sarsa Algorithm




Every **time-step**:

Policy evaluation **Sarsa**, $Q \approx q_\pi$

Policy improvement ϵ -greedy policy improvement

Q-Learning




One-step Q-learning:

$$\hat{q}(s_t, a_t) \leftarrow \hat{q}(s_t, a_t) + \alpha \left[r_{t+1} + \gamma \max_a \hat{q}(s_{t+1}, a) - \hat{q}(s_t, a_t) \right]$$

Q-Learning

$$Q^*(s, a) = E \left\{ r_{t+1} + \gamma \max_{a'} Q^*(s_{t+1}, a') \mid s_t = s, a_t = a \right\}$$



One-step Q-learning:

$$\hat{q}(s_t, a_t) \leftarrow \hat{q}(s_t, a_t) + \alpha \left[r_{t+1} + \gamma \max_a \hat{q}(s_{t+1}, a) - \hat{q}(s_t, a_t) \right]$$

Stochastic Averaging Rule

$$E(x) \approx \frac{1}{n} \sum_{i=1}^n x_i$$

Let \bar{x}_n be the average of the first n samples

$$\begin{aligned}\bar{x}_{n+1} &= \frac{1}{n+1} (x_{n+1} + n\bar{x}_n) \\ &= \frac{1}{n+1} (x_{n+1} + n\bar{x}_n + \bar{x}_n - \bar{x}_n) \\ &= \frac{1}{n+1} ((n+1)\bar{x}_n + (x_{n+1} - \bar{x}_n)) \\ &= \bar{x}_n + \frac{1}{n+1} (x_{n+1} - \bar{x}_n) \\ &= \bar{x}_n + \alpha (x_{n+1} - \bar{x}_n)\end{aligned}$$

new estimate = old estimate + α (new sample - old estimate)

Stochastic Averaging Rule

$$E(x) \approx \frac{1}{n} \sum_{i=1}^n x_i$$

Let \bar{x}_n be the average of the first n samples


$$\begin{aligned}\bar{x}_{n+1} &= \frac{1}{n+1} (x_{n+1} + n\bar{x}_n) \\ &= \frac{1}{n+1} (x_{n+1} + n\bar{x}_n + \bar{x}_n - \bar{x}_n) \\ &= \frac{1}{n+1} ((n+1)\bar{x}_n + (x_{n+1} - \bar{x}_n)) \\ &= \bar{x}_n + \frac{1}{n+1} (x_{n+1} - \bar{x}_n) \\ &= \bar{x}_n + \alpha (x_{n+1} - \bar{x}_n)\end{aligned}$$

new estimate = old estimate + α (new sample - old estimate)

$$\hat{q}(s_t, a_t) \leftarrow \hat{q}(s_t, a_t) + \alpha \left[r_{t+1} + \gamma \max_a \hat{q}(s_{t+1}, a) - \hat{q}(s_t, a_t) \right]$$

Q-Learning

$$Q^*(s, a) = E \left\{ r_{t+1} + \gamma \max_{a'} Q^*(s_{t+1}, a') \mid s_t = s, a_t = a \right\}$$




One-step Q-learning:

$$\hat{q}(s_t, a_t) \leftarrow \hat{q}(s_t, a_t) + \alpha \left[r_{t+1} + \gamma \max_a \hat{q}(s_{t+1}, a) - \hat{q}(s_t, a_t) \right]$$

Q-Learning

$$Q^*(s, a) = E \left\{ r_{t+1} + \gamma \max_{a'} Q^*(s_{t+1}, a') \mid s_t = s, a_t = a \right\}$$



One-step Q-learning:

$$\hat{q}(s_t, a_t) \leftarrow \hat{q}(s_t, a_t) + \alpha \left[r_{t+1} + \gamma \max_a \hat{q}(s_{t+1}, a) - \hat{q}(s_t, a_t) \right]$$

Temporal
Difference

Q-Learning: Off-Policy TD Control

- In off-policy control, we have two policies:
 - the behavior policy – used to generate behavior
 - estimation policy – the policy that is being evaluated and improved.

Q-learning

One-step Q-learning:

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha \left[r_{t+1} + \gamma \max_a Q(s_{t+1}, a) - Q(s_t, a_t) \right]$$

Q-learning Algorithm

Initialize $Q(s, a)$ arbitrarily

Repeat (for each episode):

 Initialize s

 Repeat (for each step of episode):

 Choose a from s using policy derived from Q (e.g., ϵ - greedy)

 Take action a , observe r, s'

$$Q(s, a) \leftarrow Q(s, a) + \alpha [r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$$

$s \leftarrow s'$;

until s is terminal

Q-learning Algorithm

Initialize $Q(s, a)$ arbitrarily

Repeat (for each episode):

Initialize s

Repeat (for each step of episode):

Choose a from s using policy derived from Q (~~e.g., ϵ -greedy~~)

Take action a , observe r, s'

$$Q(s, a) \leftarrow Q(s, a) + \alpha [r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$$

$s \leftarrow s'$;

until s is terminal

Q-learning Algorithm

Initialize $Q(s, a)$ arbitrarily

Repeat (for each episode):

Initialize s

Repeat (for each step of episode):

Choose a from s ~~using policy derived from Q (e.g., ϵ - greedy)~~


Take action a , observe r, s'

$$Q(s, a) \leftarrow Q(s, a) + \alpha [r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$$

$s \leftarrow s'$;

until s is terminal

Q-learning Algorithm



$$Q(s,a) \leftarrow Q(s,a) + \alpha[r + \gamma \max_{a'} Q(s',a') - Q(s,a)]$$

Cliff Walking: SARSA vs Q-learning

