

RAG using the llamaindex pinecone and mixtral

1 Download the Dependencies

```
[1]: !pip -q install llama-index-llms-huggingface llama-index-  
      llama-index-vector-stores-pinecone llama-index-embeddings-huggingface  
      llama-index-embeddings-instructor
```

	15.4/15.4 MB
54.7 MB/s eta 0:00:00	
	15.4/15.4 MB
29.5 MB/s eta 0:00:00	
	2.0/2.0 MB
31.8 MB/s eta 0:00:00	
	75.6/75.6 kB
5.2 MB/s eta 0:00:00	
	141.9/141.9
kB 4.8 MB/s eta 0:00:00	
	325.2/325.2
kB 16.2 MB/s eta 0:00:00	
	1.1 / 1.1 MB
57.4 MB/s eta 0:00:00	
	215.9/215.9
kB 22.2 MB/s eta 0:00:00	
	171.5/171.5
kB 10.0 MB/s eta 0:00:00	
	853.2/853.2
kB 49.6 MB/s eta 0:00:00	
	290.4/290.4
kB 28.1 MB/s eta 0:00:00	
	21.3/21.3 MB
28.4 MB/s eta 0:00:00	
	302.6/302.6
kB 33.5 MB/s eta 0:00:00	
	77.9/77.9 kB
9.9 MB/s eta 0:00:00	

7.2 MB/s eta 0:00:00 58.3/58.3 kB
6.5 MB/s eta 0:00:00 49.2/49.2 kB

2 stages of RAG

There are five key stages within RAG, which in turn will be a part of any larger application you build. These are:

Loading: this refers to getting your data from where it lives – whether it's text files, PDFs, another website, a database, or an API – into your pipeline. LlamaHub provides hundreds of connectors to choose from.

Indexing: this means creating a data structure that allows for querying the data. For LLMs this nearly always means creating vector embeddings, numerical representations of the meaning of your data, as well as numerous other metadata strategies to make it easy to accurately find contextually relevant data.

Storing: once your data is indexed you will almost always want to store your index, as well as other metadata, to avoid having to re-index it.

Querying: for any given indexing strategy there are many ways you can utilize LLMs and LlamaIndex data structures to query, including sub-queries, multi-step queries and hybrid strategies.

Evaluation: a critical step in any pipeline is checking how effective it is relative to other strategies, or when you make changes. Evaluation provides objective measures of how accurate, faithful and fast your responses to queries

```
[2]: from llama_index.core import VectorStoreIndex, SimpleDirectoryReader
```

```
[ ]: !mkdir data
```

```
[3]: loader=SimpleDirectoryReader('/content/data')  
documents=loader.load_data()
```

```
[ ]: documents
```

3 Define the embedding model

```
[5]: from llama_index.embeddings.huggingface import HuggingFaceEmbedding  
embed_model=HuggingFaceEmbedding(model_name="sentence-transformers/  
all-mpnet-base-v2")
```

/usr/local/lib/python3.10/dist-packages/huggingface_hub/utils/_token.py:89:

UserWarning:

The secret `HF_TOKEN` does not exist in your Colab secrets.

To authenticate with the Hugging Face Hub, create a token in your settings tab

(<https://huggingface.co/settings/tokens>), set it as secret in your Google Colab and restart your session.

You will be able to reuse this secret in all of your notebooks.

Please note that authentication is recommended but still optional to access public models or datasets.

```
warnings.warn(
modules.json: 0%|          | 0.00/349 [00:00<?, ?B/s]
config_sentence_transformers.json: 0%|          | 0.00/116 [00:00<?, ?B/s]
README.md: 0%|| 0.00/10.6k [00:00<?, ?B/s]
sentence_bert_config.json: 0%|          | 0.00/53.0 [00:00<?, ?B/s]
/usr/local/lib/python3.10/dist-packages/huggingface_hub/file_download.py:1132:
FutureWarning: `resume_download` is deprecated and will be removed in version
1.0.0. Downloads always resume when possible. If you want to force a new
download, use `force_download=True`.
```

```
warnings.warn(
config.json: 0%|          | 0.00/571 [00:00<?, ?B/s]
model.safetensors: 0%|          | 0.00/438M [00:00<?, ?B/s]
tokenizer_config.json: 0%|          | 0.00/363 [00:00<?, ?B/s]
vocab.txt: 0%|          | 0.00/232k [00:00<?, ?B/s]
tokenizer.json: 0%|          | 0.00/466k [00:00<?, ?B/s]
special_tokens_map.json: 0%|          | 0.00/239 [00:00<?, ?B/s]
1_Pooling/config.json: 0%|          | 0.00/190 [00:00<?, ?B/s]
```

4 Define the Prompt

```
[6]: from llama_index.core import PromptTemplate
```

```
[7]: from llama_index.llms.huggingface import HuggingFaceLLM
```

```
/usr/local/lib/python3.10/dist-packages/pydantic/_internal/_fields.py:160:
UserWarning: Field "model_id" has conflict with protected namespace "model_".
```

You may be able to resolve this warning by setting

```
`model_config['protected_namespaces'] = ()`.
```

```
warnings.warn(
```

```
[8]: system_prompt="""<|SYSTEM|># You are a Q&A assistant. Your goal is to answer_
      questions as
      accurately as possible based on the instructions and context provided
      """
```

```
[9]: # This will wrap the default prompts that are internal to llama-index
query_wrapper_prompt = PromptTemplate("<|USER|>{query_str}<|ASSISTANT|>")
```

5 Define the llm model

```
[10]: !huggingface-cli login
```

```

      _|      _|  _|      _|      _|_|_|      _|_|_|      _|_|_|      _|      _|      _|_|_|
_|_|_|_|      _|_|      _|_|_|      _|_|_|_|
      _|      _|  _|      _|  _|      _|      _|      _|_|      _|      _|
_|      _|  _|      _|      _|      _|      _|      _|      _|      _|      _|
      _|_|_|_|      _|      _|      _|      _|      _|      _|      _|      _|      _|
_|_|_|      _|_|_|_|      _|      _|_|_|_|
      _|      _|      _|      _|      _|      _|      _|      _|      _|_|      _|      _|
_|      _|      _|      _|      _|      _|      _|      _|      _|      _|_|      _|
      _|      _|      _|_|      _|_|_|      _|_|_|      _|_|_|      _|      _|      _|_|_|      _|
_|      _|      _|_|_|      _|_|_|_|

```

To login, `huggingface_hub` requires a token generated from <https://huggingface.co/settings/tokens>.

Enter your token (input will not be visible):

Add token as git credential? (Y/n) y

Token is valid (permission: write).

Cannot authenticate through git-credential as no helper is defined on your machine.

You might have to re-authenticate when pushing to the Hugging Face Hub.

Run the following command in your terminal in case you want to set the 'store' credential helper as default.

git config --global credential.helper store

Read <https://git-scm.com/book/en/v2/Git-Tools-Credential-Storage> for more details.

Token has not been saved to git credential helper.

Your token has been saved to /root/.cache/huggingface/token

Login successful

```
[11]: import torch
llm=HuggingFaceLLM(
    context_window=4096,
    max_new_tokens=256,
    generate_kwargs={"temperature": 0.7, "do_sample": False},
    system_prompt=system_prompt,
    query_wrapper_prompt=query_wrapper_prompt,
```

```

tokenizer_name="mistralai/Mistral-7B-Instruct-v0.1",
model_name="mistralai/Mistral-7B-Instruct-v0.1",
device_map="auto",
stopping_ids=[50278, 50279, 50277, 1, 0],
tokenizer_kwargs={"max_length": 4096},
# uncomment this if using CUDA to reduce memory usage
model_kwargs={"torch_dtype": torch.float16}
)

```

```

config.json: 0%|          | 0.00/571 [00:00<?, ?B/s]
model.safetensors.index.json: 0%| | 0.00/25.1k [00:00<?, ?B/s]
Downloading shards: 0%|          | 0/2 [00:00<?, ?it/s]
  model-00001-of-00002.safetensors: 0%|          | 0.00/9.94G [00:00<?, ?B/s]
  model-00002-of-00002.safetensors: 0%|          | 0.00/4.54G [00:00<?, ?B/s]
Loading checkpoint shards: 0%|          | 0/2 [00:00<?, ?it/s]
generation_config.json: 0%| | 0.00/116 [00:00<?, ?B/s]
WARNING:root:Some parameters are on the meta device device because they were
offloaded to the cpu.
tokenizer_config.json: 0%|          | 0.00/1.47k [00:00<?, ?B/s]
tokenizer.model: 0%|          | 0.00/493k [00:00<?, ?B/s]
tokenizer.json: 0%|          | 0.00/1.80M [00:00<?, ?B/s]
special_tokens_map.json: 0%|          | 0.00/72.0 [00:00<?, ?B/s]

```

6 Indexing the data

```
[12]: from llama_index.core import ServiceContext
```

```

Service_Context=ServiceContext.from_defaults(
    chunk_size=1024,
    llm=llm,
    embed_model=embed_model,
    chunk_overlap=120
)

```

```

<ipython-input-12-b3e0f15b44f2>:3: DeprecationWarning: Call to deprecated class
method from_defaults. (ServiceContext is deprecated, please use
`llama_index.settings.Settings` instead.) -- Deprecated since version 0.10.0.
  Service_Context=ServiceContext.from_defaults(

```

```
[13]: from llama_index.core import VectorStoreIndex
```

```
[14]: index=VectorStoreIndex.from_documents(documents,service_context=Service_Context)
```

```
[15]: query_engine=index.as_query_engine(k=1)
```

```
[16]: result=query_engine.query("what is Peft ?")
```

```
/usr/local/lib/python3.10/dist-  
packages/transformers/generation/configuration_utils.py:515: UserWarning:  
`do_sample` is set to `False`. However, `temperature` is set to `0.7` -- this  
flag is only used in sample-based generation modes. You should set  
`do_sample=True` or unset `temperature`.  
warnings.warn(  
Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.
```

```
[17]: result.response
```

[17] : 'PeFT stands for Parameter Efficient Fine-tuning. It is a technique used in natural language processing (NLP) to fine-tune pre-trained language models (PLMs) on a smaller dataset. PeFT allows for selective or not updating pretrained parameters, which helps to avoid overfitting. There has been a surge in interest in PeFT methods, and several surveys have been conducted on the topic. However, existing surveys have limitations, such as not covering much of the latest work in the field or not performing relevant experiments. In this work, the authors aim to provide a comprehensive and systematic study of PeFT methods for PLMs in NLP, including categorizing the methods, providing detailed explanations of the ideas and specific implementations, comparing the similarities and differences among various types of PeFT methods, and conducting extensive fine-tuning experiments with 11 representative PeFT methods.'

```
[18] : print(result)
```

PeFT stands for Parameter Efficient Fine-tuning. It is a technique used in natural language processing (NLP) to fine-tune pre-trained language models (PLMs) on a smaller dataset. PeFT allows for selective or not updating pretrained parameters, which helps to avoid overfitting. There has been a surge in interest in PeFT methods, and several surveys have been conducted on the topic. However, existing surveys have limitations, such as not covering much of the latest work in the field or not performing relevant experiments. In this work, the authors aim to provide a comprehensive and systematic study of PeFT methods for PLMs in NLP, including categorizing the methods, providing detailed explanations of the ideas and specific implementations, comparing the similarities and differences among various types of PeFT methods, and conducting extensive fine-tuning experiments with 11 representative PeFT methods.

```
[19] : from IPython.display import Markdown, display
```

```
[20] : display(Markdown(f"<b>{result}</b>"))
```

PeFT stands for Parameter Efficient Fine-tuning. It is a technique used in natural language processing (NLP) to fine-tune pre-trained language models (PLMs) on a smaller dataset. PeFT allows for selective or not updating pretrained parameters, which helps to avoid overfitting. There has been a surge in interest in PeFT methods, and several surveys have been conducted on the topic. However, existing surveys have limitations, such as not covering much of the latest work in the field or not performing relevant experiments. In this work, the authors aim to provide a comprehensive and systematic study of PeFT methods for PLMs in NLP, including categorizing the methods, providing detailed explanations of the ideas and specific implementations, comparing the similarities and differences among various types of PeFT methods, and conducting extensive fine-tuning experiments with 11 representative PeFT methods.

7 Storing the data

```
[21] : # fixing unicode error in google colab
import locale
locale.getpreferredencoding = lambda: "UTF-8"

[22] : import os
import pinecone

PINECONE_API_KEY = os.environ.get('PINECONE_API_KEY',
                                   '48180f5d-52d0-49e6-b8b5-89638588ed95')
PINECONE_API_ENV = os.environ.get('PINECONE_API_ENV', 'gcp-starter')

# Create a Pinecone instance for initialization
pinecone_client = pinecone.Pinecone(
    api_key=PINECONE_API_KEY,
    environment=PINECONE_API_ENV
)

# Now you can use pinecone_client to interact with Pinecone
# For example, to list existing indexes:
existing_indexes = pinecone_client.list_indexes()
print(existing_indexes)

{'indexes': [{'dimension': 768,
               'host': 'rushy-5xyrzmy.svc.aped-4627-b74a.pinecone.io',
               'metric': 'cosine',
               'name': 'rushy',
               'spec': {'serverless': {'cloud': 'aws', 'region': 'us-east-1'}},
               'status': {'ready': True, 'state': 'Ready'}}]}

[23] : pinecone_index = pinecone_client.Index("rushy")

[24] : from llama_index.vector_stores.pinecone import PineconeVectorStore
from llama_index.core import VectorStoreIndex
from llama_index.core import StorageContext
```

```
[25] : vector_store = PineconeVectorStore(  
        pinecone_index=pinecone_index, namespace="test_05_14"  
    )
```

```
[26] : storage_context = StorageContext.from_defaults(vector_store=vector_store)
```

```
[27] : index = VectorStoreIndex.from_documents(  
        documents, storage_context=storage_context, embed_model=embed_model  
    )
```

Upserted vectors: 0% | 0/47 [00:00<?, ?it/s]

8 Querying the data

```
[28] : query_engine = index.as_query_engine(llm=llm)
```

```
[29] : result3=query_engine.query("what is peft?")
```

```
/usr/local/lib/python3.10/dist-  
packages/transformers/generation/configuration_utils.py:515: UserWarning:  
`do_sample` is set to `False`. However, `temperature` is set to `0.7` -- this  
flag is only used in sample-based generation modes. You should set  
`do_sample=True` or unset `temperature`.  
    warnings.warn(  
Setting `pad_token_id` to `eos_token_id`:2 for open-end generation.
```

```
[30] : result3.response
```

[30]: 'PEFT stands for "Parameter-Efficient Fine-Tuning". It is a technique used to fine-tune pre-trained language models on specific tasks with fewer parameters than the original model. This is achieved by using a combination of techniques such as pruning, quantization, and knowledge distillation. PEFT methods can significantly reduce the number of trainable parameters while maintaining or improving the performance on the task.'